

shitij-HQLSTM-Perplexity-Priority-1-Better-5

October 26, 2025

[]:

```
[42]: import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import pennylane as qml

# =====
# REPRODUCIBILITY / SEED
# =====
SEED = 42
import random
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False

# =====
# DATA PREPARATION
# =====
# Paths
train_path = r"C:\Users\acer\OneDrive\QPoland\X_train.csv"
test_path = r"C:\Users\acer\OneDrive\QPoland\X_test.csv"

X_train = pd.read_csv(train_path).sort_values("Date")
X_test = pd.read_csv(test_path).sort_values("Date")

def add_simple_features(df):
```

```

df = df.copy()
df['Price_Range'] = df['High'] - df['Low']
df['HL_Ratio'] = df['High'] / (df['Low'] + 1e-10)
df['Volume_MA_3'] = df['Volume'].rolling(window=3, min_periods=1).mean()
return df.bfill().ffill()

X_train = add_simple_features(X_train)
X_test = add_simple_features(X_test)

features = ["Open", "High", "Low", "Volume", "Price_Range", "HL_Ratio",
            ↪ "Volume_MA_3"]
target = "Close"

scaler_X = MinMaxScaler()
scaler_y = MinMaxScaler()
X_train_scaled = scaler_X.fit_transform(X_train[features])
y_train_scaled = scaler_y.fit_transform(X_train[[target]])

SEQ_LEN = 8

def create_sequences(X, y, seq_len):
    X_seq, y_seq = [], []
    for i in range(len(X) - seq_len):
        X_seq.append(X[i:i+seq_len])
        y_seq.append(y[i+seq_len])
    return np.array(X_seq), np.array(y_seq)

X_seq, y_seq = create_sequences(X_train_scaled, y_train_scaled, SEQ_LEN)

split = int(len(X_seq) * 0.65)
X_train_t = torch.tensor(X_seq[:split], dtype=torch.float32)
y_train_t = torch.tensor(y_seq[:split], dtype=torch.float32)
X_val_t = torch.tensor(X_seq[split:], dtype=torch.float32)
y_val_t = torch.tensor(y_seq[split:], dtype=torch.float32)

print(f"Training: {len(X_train_t)}, Validation: {len(X_val_t)}, Seq Length:
    ↪ {SEQ_LEN}\n")

# =====
# HYBRID LSTM + QNN PARAMETERS
# =====
input_size = len(features)
n_qubits = 4 # max qubits
hidden_size = n_qubits
dropout_rate = 0.1 # CHANGED: Increased from 0.05 to 0.2 to combat overfitting
n_layers = 2

```

```

def xavier_init(in_size, out_size):
    limit = np.sqrt(6.0 / (in_size + out_size))
    return torch.randn(in_size, out_size) * limit

# Classical LSTM gates
W_ii = nn.Parameter(xavier_init(input_size, hidden_size))
W_hi = nn.Parameter(xavier_init(hidden_size, hidden_size))
b_i = nn.Parameter(torch.zeros(hidden_size))

W_if = nn.Parameter(xavier_init(input_size, hidden_size))
W_hf = nn.Parameter(xavier_init(hidden_size, hidden_size))
b_f = nn.Parameter(torch.ones(hidden_size))

W_io = nn.Parameter(xavier_init(input_size, hidden_size))
W_ho = nn.Parameter(xavier_init(hidden_size, hidden_size))
b_o = nn.Parameter(torch.zeros(hidden_size))

W_out = nn.Parameter(xavier_init(hidden_size, 1))
b_out = nn.Parameter(torch.zeros(1))

dropout = nn.Dropout(dropout_rate)

# =====
# QUANTUM LAYER SETUP
# =====
dev = qml.device("default.qubit", wires=n_qubits)

shape = qml.BasicEntanglerLayers.shape(n_layers=n_layers, n_wires=n_qubits)
weights = nn.Parameter(0.01 * torch.randn(*shape))

# @qml.qnode(dev, interface="torch")
# def quantum_node(inputs, weights):
#     qml.templates.AngleEmbedding(inputs, wires=range(n_qubits))
#     qml.BasicEntanglerLayers(weights=weights, wires=range(n_qubits))
#     return [qml.expval(qml.PauliZ(i)) for i in range(n_qubits)]

@qml.qnode(dev, interface="torch", diff_method="backprop",
            max_diff_method="finite-diff")
def quantum_node_batch(inputs, weights):
    # inputs shape: (batch_size, n_qubits)
    qml.templates.AngleEmbedding(inputs, wires=range(n_qubits))
    qml.BasicEntanglerLayers(weights=weights, wires=range(n_qubits))
    return [qml.expval(qml.PauliZ(i)) for i in range(n_qubits)]

```

```

# PCA transformer to reduce input to 4 qubits
pca = PCA(n_components=n_qubits)

# Collect parameters
lstm_params = [W_ii, W_hi, b_i, W_if, W_hf, b_f, W_io, W_ho, b_o, W_out, b_out,
               ↪weights]

# =====
# HYBRID LSTM FORWARD FUNCTION
# =====
def lstm_forward(x_batch, training=True):
    batch_size = x_batch.shape[0]
    seq_length = x_batch.shape[1]
    h_t = torch.zeros(batch_size, hidden_size)
    c_t = torch.zeros(batch_size, hidden_size)

    for t in range(seq_length):
        x_t = x_batch[:, t, :]
        i_t = torch.sigmoid(x_t @ W_ii + h_t @ W_hi + b_i)
        f_t = torch.sigmoid(x_t @ W_if + h_t @ W_hf + b_f)
        o_t = torch.sigmoid(x_t @ W_io + h_t @ W_ho + b_o)

        # Quantum cell candidate
        concat_input = torch.cat([x_t, h_t], dim=1)
        q_input = torch.tensor(pca.fit_transform(concat_input.detach()).
                               ↪numpy()), dtype=torch.float32)

        q_outputs = []
        for i in range(q_input.shape[0]):
            q_out = quantum_node(q_input[i], weights)
            q_outputs.append(torch.tensor(q_out, dtype=torch.float32))
        c_tilde = torch.stack(q_outputs, dim=0)
        c_tilde = torch.tanh(c_tilde)

        if training:
            c_tilde = dropout(c_tilde)

        c_t = f_t * c_t + i_t * c_tilde
        h_t = o_t * torch.tanh(c_t)

    if training:
        h_t = dropout(h_t)

    output = h_t @ W_out + b_out
    return output

```

```

# =====
# TRAINING LOOP WITH PRIORITY 1 IMPROVEMENTS
# =====
# CHANGED: Using Adam optimizer with weight decay (L2 regularization)
lr = 0.001
wd = 1e-6
optimizer = optim.Adam(lstm_params, lr=lr, weight_decay=wd)

# CHANGED: Using Huber loss instead of MSE for robustness
beta = 1.0
criterion = nn.SmoothL1Loss(beta=beta)

EPOCHS = 1000 # CHANGED: Increased max epochs (early stopping will prevent_
↳ overtraining)
train_losses, val_losses = [], []

# NEW: Early stopping parameters
best_val_loss = float('inf')
patience = 50 # Stop if no improvement for 10 epochs
patience_counter = 0
best_state_dict = None

print("Training Hybrid LSTM + QNN with Priority 1 improvements...")
print(f"- Dropout: {dropout_rate}")
print(f"- Learning Rate: {lr}")
print(f"- Weight Decay: {wd}")
#print(f"- Optimizer: Adam with weight_decay=1e-4")
print(f"- Loss: Huber (SmoothL1) with Beta : {beta}")
print(f"- Early stopping patience: {patience}\n")

for epoch in range(1, EPOCHS + 1):
    # Training phase
    dropout.train()
    optimizer.zero_grad()
    train_pred = lstm_forward(X_train_t, training=True)
    train_loss = criterion(train_pred, y_train_t)
    train_loss.backward()
    torch.nn.utils.clip_grad_norm_(lstm_params, max_norm=1.0)
    optimizer.step()
    train_losses.append(train_loss.item())

    # Validation phase
    dropout.eval()
    with torch.no_grad():
        val_pred = lstm_forward(X_val_t, training=False)
        val_loss = criterion(val_pred, y_val_t)

```

```

        val_losses.append(val_loss.item())

# NEW: Early stopping logic
if val_loss.item() < best_val_loss:
    best_val_loss = val_loss.item()
    patience_counter = 0
    # Save best model state
    best_state_dict = {
        'W_ii': W_ii.data.clone(),
        'W_hi': W_hi.data.clone(),
        'b_i': b_i.data.clone(),
        'W_if': W_if.data.clone(),
        'W_hf': W_hf.data.clone(),
        'b_f': b_f.data.clone(),
        'W_io': W_io.data.clone(),
        'W_ho': W_ho.data.clone(),
        'b_o': b_o.data.clone(),
        'W_out': W_out.data.clone(),
        'b_out': b_out.data.clone(),
        'weights': weights.data.clone()
    }
    print(f"Epoch {epoch:2d} | Train Loss: {train_loss.item():.6f} | Val_
↳ Loss: {val_loss.item():.6f} (best)")
else:
    patience_counter += 1
    print(f"Epoch {epoch:2d} | Train Loss: {train_loss.item():.6f} | Val_
↳ Loss: {val_loss.item():.6f} | Patience: {patience_counter}/{patience}")

# NEW: Stop training if no improvement
if patience_counter >= patience:
    print(f"\n Early stopping triggered at epoch {epoch}. Best val loss:
↳ {best_val_loss:.6f}")
    break

# NEW: Restore best model weights
if best_state_dict is not None:
    print("\n Restoring best model weights from validation...")
    W_ii.data = best_state_dict['W_ii']
    W_hi.data = best_state_dict['W_hi']
    b_i.data = best_state_dict['b_i']
    W_if.data = best_state_dict['W_if']
    W_hf.data = best_state_dict['W_hf']
    b_f.data = best_state_dict['b_f']
    W_io.data = best_state_dict['W_io']
    W_ho.data = best_state_dict['W_ho']
    b_o.data = best_state_dict['b_o']
    W_out.data = best_state_dict['W_out']

```

```

b_out.data = best_state_dict['b_out']
weights.data = best_state_dict['weights']

# =====
# PLOTTING TRAINING CURVES
# =====
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Train Loss', linewidth=2)
plt.plot(val_losses, label='Validation Loss', linewidth=2)
plt.axvline(x=len(val_losses) - patience_counter, color='red', linestyle='--',
            label=f'Best Model (epoch {len(val_losses) - patience_counter})',
            alpha=0.7)
plt.title('Hybrid LSTM + QNN Training Progress (with Early Stopping)')
plt.xlabel('Epoch')
plt.ylabel('Huber Loss')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

# =====
# EVALUATION ON TRAIN/VALIDATION
# =====
dropout.eval()
with torch.no_grad():
    train_pred_final = lstm_forward(X_train_t, training=False).numpy()
    val_pred_final = lstm_forward(X_val_t, training=False).numpy()

train_pred_rescaled = scaler_y.inverse_transform(train_pred_final)
val_pred_rescaled = scaler_y.inverse_transform(val_pred_final)
y_train_orig = scaler_y.inverse_transform(y_train_t.numpy())
y_val_orig = scaler_y.inverse_transform(y_val_t.numpy())

train_mse = mean_squared_error(y_train_orig, train_pred_rescaled)
train_r2 = r2_score(y_train_orig, train_pred_rescaled)
val_mse = mean_squared_error(y_val_orig, val_pred_rescaled)
val_r2 = r2_score(y_val_orig, val_pred_rescaled)

print("\n" + "="*60)
print("FINAL METRICS (Best Model from Early Stopping)")
print("="*60)
print(f"TRAIN - MSE: {train_mse:.4f}, R2: {train_r2:.4f}, RMSE: {np.
    sqrt(train_mse):.2f}")
print(f"VAL - MSE: {val_mse:.4f}, R2: {val_r2:.4f}, RMSE: {np.sqrt(val_mse):.
    2f}")

```

```

print("="*60)

# NEW: Show overfitting gap
r2_gap = train_r2 - val_r2
if r2_gap > 0.2:
    print(f" WARNING: Large overfitting gap detected (Train R2 - Val R2 = {r2_gap:.4f})")
    print("    Consider: increasing dropout, reducing model complexity, or adding more data")
elif r2_gap > 0.1:
    print(f" Moderate overfitting detected (Train R2 - Val R2 = {r2_gap:.4f})")
else:
    print(f" Good generalization (Train R2 - Val R2 = {r2_gap:.4f})")
print("="*60)

# =====
# PLOT VALIDATION ACTUAL vs PREDICTED
# =====
plt.figure(figsize=(12, 6))
plt.plot(y_val_orig, label='Actual', linewidth=2, marker='o', markersize=4)
plt.plot(val_pred_rescaled, label='Predicted', linewidth=2, marker='s', markersize=4)
plt.title(f'Validation: Actual vs Predicted (R2 = {val_r2:.4f})')
plt.xlabel('Sample')
plt.ylabel('Close Price')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

# =====
# PREDICTIONS ON TEST DATA
# =====
X_test_scaled = scaler_X.transform(X_test[features])
combined = np.concatenate([X_train_scaled[-SEQ_LEN:], X_test_scaled])
X_test_seq = [combined[i:i+SEQ_LEN] for i in range(len(combined) - SEQ_LEN)]
X_test_t = torch.tensor(np.array(X_test_seq), dtype=torch.float32)

with torch.no_grad():
    test_pred = lstm_forward(X_test_t, training=False).numpy()
test_pred_rescaled = scaler_y.inverse_transform(test_pred)

predicted_closes = pd.DataFrame({
    "Date": X_test["Date"].values[:len(test_pred_rescaled)],
    "Predicted_Close": test_pred_rescaled.flatten()
})

```



```

})

print("\n" + "="*60)
print("TEST SET PREDICTIONS")
print("="*60)
print(predicted_closes.to_string(index=False))
print("="*60)

# =====
# PLOT TEST PREDICTIONS
# =====
plt.figure(figsize=(12, 6))
plt.plot(X_test["Date"].values[:len(test_pred_rescaled)], test_pred_rescaled.
        ↪flatten(),
        label='Predicted', linewidth=2, marker='s', markersize=4)
plt.title('Test Set Predictions (Best Model)')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.xticks(rotation=45)
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

Training: 102, Validation: 55, Seq Length: 8

Training Hybrid LSTM + QNN with Priority 1 improvements...

- Dropout: 0.1
- Learning Rate: 0.001
- Weight Decay: 1e-06
- Loss: Huber (SmoothL1) with Beta : 1.0
- Early stopping patience: 50

C:\Users\acer\anaconda3\Lib\site-packages\pennyLane\workflow\qnode.py:135:

UserWarning: Received gradient_kwarg max_diff_method, which is not included in the list of standard qnode gradient kwargs.

warnings.warn(


Epoch	1		Train Loss: 0.233422		Val Loss: 0.055734	(best)
Epoch	2		Train Loss: 0.230053		Val Loss: 0.053522	(best)
Epoch	3		Train Loss: 0.245483		Val Loss: 0.051355	(best)
Epoch	4		Train Loss: 0.212411		Val Loss: 0.049249	(best)
Epoch	5		Train Loss: 0.208945		Val Loss: 0.047194	(best)
Epoch	6		Train Loss: 0.212351		Val Loss: 0.045192	(best)
Epoch	7		Train Loss: 0.212122		Val Loss: 0.043229	(best)
Epoch	8		Train Loss: 0.210388		Val Loss: 0.041314	(best)
Epoch	9		Train Loss: 0.197015		Val Loss: 0.039460	(best)

Epoch 10	Train Loss: 0.209175	Val Loss: 0.037645	(best)
Epoch 11	Train Loss: 0.200148	Val Loss: 0.035876	(best)
Epoch 12	Train Loss: 0.177837	Val Loss: 0.034163	(best)
Epoch 13	Train Loss: 0.176740	Val Loss: 0.032503	(best)
Epoch 14	Train Loss: 0.186339	Val Loss: 0.030900	(best)
Epoch 15	Train Loss: 0.181005	Val Loss: 0.029352	(best)
Epoch 16	Train Loss: 0.184694	Val Loss: 0.027845	(best)
Epoch 17	Train Loss: 0.178018	Val Loss: 0.026384	(best)
Epoch 18	Train Loss: 0.164018	Val Loss: 0.024976	(best)
Epoch 19	Train Loss: 0.162167	Val Loss: 0.023631	(best)
Epoch 20	Train Loss: 0.160975	Val Loss: 0.022328	(best)
Epoch 21	Train Loss: 0.166571	Val Loss: 0.021076	(best)
Epoch 22	Train Loss: 0.148871	Val Loss: 0.019870	(best)
Epoch 23	Train Loss: 0.156160	Val Loss: 0.018712	(best)
Epoch 24	Train Loss: 0.149582	Val Loss: 0.017604	(best)
Epoch 25	Train Loss: 0.151916	Val Loss: 0.016549	(best)
Epoch 26	Train Loss: 0.143517	Val Loss: 0.015535	(best)
Epoch 27	Train Loss: 0.138224	Val Loss: 0.014570	(best)
Epoch 28	Train Loss: 0.123705	Val Loss: 0.013651	(best)
Epoch 29	Train Loss: 0.133458	Val Loss: 0.012779	(best)
Epoch 30	Train Loss: 0.155420	Val Loss: 0.011951	(best)
Epoch 31	Train Loss: 0.140851	Val Loss: 0.011166	(best)
Epoch 32	Train Loss: 0.129420	Val Loss: 0.010426	(best)
Epoch 33	Train Loss: 0.117455	Val Loss: 0.009731	(best)
Epoch 34	Train Loss: 0.116433	Val Loss: 0.009080	(best)
Epoch 35	Train Loss: 0.124770	Val Loss: 0.008467	(best)
Epoch 36	Train Loss: 0.127047	Val Loss: 0.007890	(best)
Epoch 37	Train Loss: 0.114883	Val Loss: 0.007354	(best)
Epoch 38	Train Loss: 0.098370	Val Loss: 0.006853	(best)
Epoch 39	Train Loss: 0.099923	Val Loss: 0.006391	(best)
Epoch 40	Train Loss: 0.112040	Val Loss: 0.005959	(best)
Epoch 41	Train Loss: 0.098198	Val Loss: 0.005559	(best)
Epoch 42	Train Loss: 0.096957	Val Loss: 0.005193	(best)
Epoch 43	Train Loss: 0.089948	Val Loss: 0.004859	(best)
Epoch 44	Train Loss: 0.105086	Val Loss: 0.004558	(best)
Epoch 45	Train Loss: 0.089044	Val Loss: 0.004287	(best)
Epoch 46	Train Loss: 0.106175	Val Loss: 0.004043	(best)
Epoch 47	Train Loss: 0.092981	Val Loss: 0.003829	(best)
Epoch 48	Train Loss: 0.091996	Val Loss: 0.003640	(best)
Epoch 49	Train Loss: 0.102632	Val Loss: 0.003476	(best)
Epoch 50	Train Loss: 0.080826	Val Loss: 0.003338	(best)
Epoch 51	Train Loss: 0.089385	Val Loss: 0.003226	(best)
Epoch 52	Train Loss: 0.076926	Val Loss: 0.003139	(best)
Epoch 53	Train Loss: 0.094700	Val Loss: 0.003077	(best)
Epoch 54	Train Loss: 0.081696	Val Loss: 0.003036	(best)
Epoch 55	Train Loss: 0.079432	Val Loss: 0.003018	(best)
Epoch 56	Train Loss: 0.080826	Val Loss: 0.003021	Patience: 1/50
Epoch 57	Train Loss: 0.084116	Val Loss: 0.003044	Patience: 2/50

Epoch 58		Train Loss: 0.063110		Val Loss: 0.003088		Patience: 3/50
Epoch 59		Train Loss: 0.081234		Val Loss: 0.003148		Patience: 4/50
Epoch 60		Train Loss: 0.089714		Val Loss: 0.003223		Patience: 5/50
Epoch 61		Train Loss: 0.064980		Val Loss: 0.003312		Patience: 6/50
Epoch 62		Train Loss: 0.071085		Val Loss: 0.003418		Patience: 7/50
Epoch 63		Train Loss: 0.070749		Val Loss: 0.003536		Patience: 8/50
Epoch 64		Train Loss: 0.075675		Val Loss: 0.003668		Patience: 9/50
Epoch 65		Train Loss: 0.068577		Val Loss: 0.003813		Patience: 10/50
Epoch 66		Train Loss: 0.057230		Val Loss: 0.003972		Patience: 11/50
Epoch 67		Train Loss: 0.058779		Val Loss: 0.004146		Patience: 12/50
Epoch 68		Train Loss: 0.059905		Val Loss: 0.004330		Patience: 13/50
Epoch 69		Train Loss: 0.067675		Val Loss: 0.004525		Patience: 14/50
Epoch 70		Train Loss: 0.040930		Val Loss: 0.004730		Patience: 15/50
Epoch 71		Train Loss: 0.068092		Val Loss: 0.004940		Patience: 16/50
Epoch 72		Train Loss: 0.060276		Val Loss: 0.005164		Patience: 17/50
Epoch 73		Train Loss: 0.075793		Val Loss: 0.005392		Patience: 18/50
Epoch 74		Train Loss: 0.052472		Val Loss: 0.005631		Patience: 19/50
Epoch 75		Train Loss: 0.064348		Val Loss: 0.005877		Patience: 20/50
Epoch 76		Train Loss: 0.061953		Val Loss: 0.006135		Patience: 21/50
Epoch 77		Train Loss: 0.078071		Val Loss: 0.006394		Patience: 22/50
Epoch 78		Train Loss: 0.055053		Val Loss: 0.006667		Patience: 23/50
Epoch 79		Train Loss: 0.054929		Val Loss: 0.006950		Patience: 24/50
Epoch 80		Train Loss: 0.045714		Val Loss: 0.007240		Patience: 25/50
Epoch 81		Train Loss: 0.055936		Val Loss: 0.007538		Patience: 26/50
Epoch 82		Train Loss: 0.075177		Val Loss: 0.007838		Patience: 27/50
Epoch 83		Train Loss: 0.039239		Val Loss: 0.008148		Patience: 28/50
Epoch 84		Train Loss: 0.046254		Val Loss: 0.008467		Patience: 29/50
Epoch 85		Train Loss: 0.074730		Val Loss: 0.008784		Patience: 30/50
Epoch 86		Train Loss: 0.042992		Val Loss: 0.009102		Patience: 31/50
Epoch 87		Train Loss: 0.048612		Val Loss: 0.009426		Patience: 32/50
Epoch 88		Train Loss: 0.061960		Val Loss: 0.009743		Patience: 33/50
Epoch 89		Train Loss: 0.058667		Val Loss: 0.010063		Patience: 34/50
Epoch 90		Train Loss: 0.052734		Val Loss: 0.010383		Patience: 35/50
Epoch 91		Train Loss: 0.045166		Val Loss: 0.010700		Patience: 36/50
Epoch 92		Train Loss: 0.058232		Val Loss: 0.011013		Patience: 37/50
Epoch 93		Train Loss: 0.065347		Val Loss: 0.011322		Patience: 38/50
Epoch 94		Train Loss: 0.072238		Val Loss: 0.011617		Patience: 39/50
Epoch 95		Train Loss: 0.048796		Val Loss: 0.011904		Patience: 40/50
Epoch 96		Train Loss: 0.050282		Val Loss: 0.012198		Patience: 41/50
Epoch 97		Train Loss: 0.060904		Val Loss: 0.012477		Patience: 42/50
Epoch 98		Train Loss: 0.036803		Val Loss: 0.012770		Patience: 43/50
Epoch 99		Train Loss: 0.049025		Val Loss: 0.013061		Patience: 44/50
Epoch 100		Train Loss: 0.074740		Val Loss: 0.013337		Patience: 45/50
Epoch 101		Train Loss: 0.040063		Val Loss: 0.013611		Patience: 46/50
Epoch 102		Train Loss: 0.055178		Val Loss: 0.013879		Patience: 47/50
Epoch 103		Train Loss: 0.056842		Val Loss: 0.014149		Patience: 48/50
Epoch 104		Train Loss: 0.042425		Val Loss: 0.014426		Patience: 49/50
Epoch 105		Train Loss: 0.051506		Val Loss: 0.014695		Patience: 50/50

Early stopping triggered at epoch 105. Best val loss: 0.003018

Restoring best model weights from validation...



=====

FINAL METRICS (Best Model from Early Stopping)

=====


TRAIN - MSE: 163798.4531, R^2 : -1.2207, RMSE: 404.72

VAL - MSE: 13927.7197, R^2 : 0.3460, RMSE: 118.02

=====

Good generalization (Train R^2 - Val R^2 = -1.5667)

=====




=====

TEST SET PREDICTIONS

=====

Date	Predicted_Close
2025-09-02	6381.488770
2025-09-03	6316.316895
2025-09-04	6370.878906
2025-09-05	6325.871094
2025-09-08	6287.756836
2025-09-09	6412.341309
2025-09-10	6375.680176
2025-09-11	6357.419922
2025-09-12	6335.991211
2025-09-15	6406.552246

=====



[]:

[]:

[]:

[44]: `import torch`

```
# Save all parameters to a dictionary
model_params = {
    'W_ii': W_ii.data,
    'W_hi': W_hi.data,
    'b_i': b_i.data,
    'W_if': W_if.data,
    'W_hf': W_hf.data,
    'b_f': b_f.data,
    'W_io': W_io.data,
    'W_ho': W_ho.data,
    'b_o': b_o.data,
    'W_out': W_out.data,
    'b_out': b_out.data,
    'weights': weights.data,
    "scaler_X": scaler_X,      # Save feature scaler
    "scaler_y": scaler_y,     # Save target scaler
    "pca": pca,               # Save PCA transformer
}

# Choose your output file name
```

```

torch.save(model_params, 'hybrid_qlstm_model_5.pth')
print("Saved model parameters to hybrid_qlstm_model_5.pth")

'''
# Load model parameters
checkpoint = torch.load('hybrid_qlstm_model.pth')

# Restore each parameter
W_ii.data = checkpoint['W_ii']
W_hi.data = checkpoint['W_hi']
b_i.data = checkpoint['b_i']
W_if.data = checkpoint['W_if']
W_hf.data = checkpoint['W_hf']
b_f.data = checkpoint['b_f']
W_io.data = checkpoint['W_io']
W_ho.data = checkpoint['W_ho']
b_o.data = checkpoint['b_o']
W_out.data = checkpoint['W_out']
b_out.data = checkpoint['b_out']
weights.data = checkpoint['weights']

# Restore scalars and PCA
scaler_X = checkpoint['scaler_X']
scaler_y = checkpoint['scaler_y']
pca = checkpoint['pca']

print("Loaded model parameters from hybrid_qlstm_model.pth")
'''

import json
def to_serializable(val):
    if isinstance(val, (np.generic, np.ndarray)):
        return val.item() # Converts NumPy scalars to native Python types
    elif isinstance(val, torch.Tensor):
        return val.item() # Converts 0-dim tensor to Python scalar
    else:
        return val

run_info = {k: to_serializable(v) for k, v in {
    "train_mse": train_mse,
    "train_r2": train_r2,
    "val_mse": val_mse,
    "val_r2": val_r2,
    "dropout_rate": dropout_rate,
    "early_stopping_patience": patience,
    "epochs_run": epoch
}.items()}
```

```
with open("hybrid_qlstm_metadata_5.json", "w") as f:
    json.dump(run_info, f, indent=2)
```

Saved model parameters to hybrid_qlstm_model_5.pth

```
[92]: import pandas as pd

# Data as a dictionary
data = {
    "Date": [
        "2025-09-02", "2025-09-03", "2025-09-04", "2025-09-05", "2025-09-08",
        "2025-09-09", "2025-09-10", "2025-09-11", "2025-09-12", "2025-09-15"
    ],
    "Predicted_Close": [
        6381.488770, 6316.316895, 6370.878906, 6325.871094, 6287.756836,
        6412.341309, 6375.680176, 6357.419922, 6335.991211, 6406.552246
    ]
}

# Create DataFrame
df = pd.DataFrame(data)

# Save to CSV file
df.to_csv("test_set_predictions.csv", index=False)

print("Data saved to 'test_set_predictions.csv'.")
```

Data saved to 'test_set_predictions.csv'.

```
[ ]:
```